

A classic locked-room mystery.
Eve was in the false branch of a
conditional the whole time,
how could she do it?

 Creative Commons Attribution-ShareAlike 4.0
Mozilla Research | DePaul University | U. California San Diego

Overview

Introduction
Model
Attacks
Experiments
Conclusions

The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

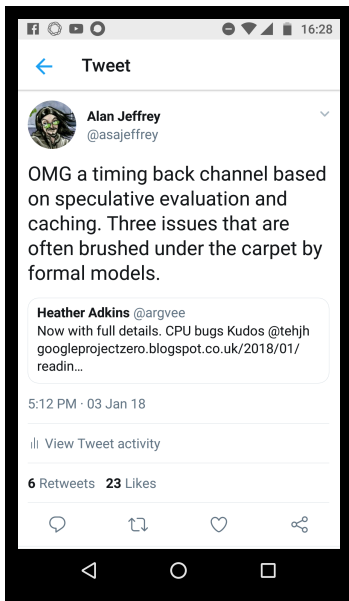
Model

Attacks

Experiments

Conclusions

Why? Spectre!



The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

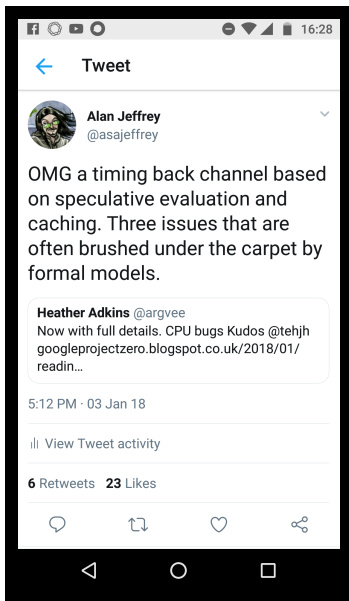
Model

Attacks

Experiments

Conclusions

Why? Spectre!



Attacks bypass dynamic security checks:

```
if (canReadSecret) {  
    doStuffWith(SECRET);  
}
```

Information flow from SECRET even though `canReadSecret` is false.

Most formal models ignore code in branches that aren't taken.

The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

Model

Attacks

Experiments

Conclusions

Models that include speculation?

The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

There are some models that include speculation
relaxed memory models:

- ▶ *The Java Memory Model*
Manson, Pugh and Adve, 2005.
- ▶ *Generative Operational Semantics for Relaxed Memory Models*
Jagadeesan, Pitcher and Riely, 2010.
- ▶ *A promising semantics for relaxed-memory concurrency*
Kang, Hur, Lahav, Vafeiadis and Dreyer, 2017.

Introduction

Model

Attacks

Experiments

Conclusions

Models that include speculation?

The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

There are some models that include speculation
relaxed memory models:

- ▶ *The Java Memory Model*
Manson, Pugh and Adve, 2005.
- ▶ *Generative Operational Semantics for Relaxed Memory Models*
Jagadeesan, Pitcher and Riely, 2010.
- ▶ *A promising semantics for relaxed-memory concurrency*
Kang, Hur, Lahav, Vafeiadis and Dreyer, 2017.

Question: is there a simple model similar to those of relaxed memory, that can model speculation?

Introduction

Model

Attacks

Experiments

Conclusions

Information flow attacks on speculation

The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Speculation happens in many places:

- ▶ *Speculation in hardware* (branch prediction, . . .)
- ▶ *Transactions* (transactional memory, . . .)
- ▶ *Relaxed memory* (compiler optimizations, . . .)

Introduction

Model

Attacks

Experiments

Conclusions

Information flow attacks on speculation

The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Speculation happens in many places:

- ▶ *Speculation in hardware* (branch prediction, . . .)
Attacked by Spectre (Kocher *et al.* 2019).
- ▶ *Transactions* (transactional memory, . . .)
Attacked by Prime+Abort (Disselkoen *et al.* 2017).
- ▶ *Relaxed memory* (compiler optimizations, . . .)
No known attacks.

Question: are there information flow attacks against compiler optimizations?

Introduction

Model

Attacks

Experiments

Conclusions

Contributions

- ▶ A simple compositional model.
- ▶ Examples.
- ▶ Attacks (including a new attack on relaxed memory).
- ▶ Experiments (testing practicality of new attacks).

The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

Model

Attacks

Experiments

Conclusions

C11-style models are based on *events*
with *labels* (e.g. $(R \times 3)$ or $(W \times 3)$)
and *relations* (e.g. happens-before or reads-from).

C11-style models are based on *events*
with *labels* (e.g. $(R \times 3)$ or $(W \times 3)$)
and *relations* (e.g. happens-before or reads-from).


Simplest such is *partially ordered multisets* (Gisher, 1988).

Only one relation, a partial order modelling dependency

C11-style models are based on *events*
with *labels* (e.g. $(R \times 3)$ or $(W \times 3)$)
and *relations* (e.g. happens-before or reads-from).

Simplest such is *partially ordered multisets* (Gisher, 1988).

Only one relation, a partial order modelling dependency, e.g.



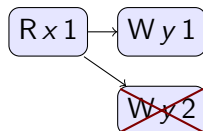
is an execution of $(r := x; y := 1; z := r + 1)$.

Pomsets

C11-style models are based on *events* with *labels* (e.g. $(R \times 3)$ or $(W \times 3)$) and *relations* (e.g. happens-before or reads-from).

Simplest such is *partially ordered multisets* (Gisher, 1988).

Only one relation, a partial order modelling dependency, e.g.



is an execution of $(\text{if } (x) \{ y := 1 \} \text{ else } \{ y := 2 \})$.

Compositional pomset model

First off, straight-line code.

The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

Model

Attacks

Experiments

Conclusions

Compositional pomset model

First off, straight-line code.

New idea: put preconditions on events

The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

Model

Attacks

Experiments

Conclusions

Compositional pomset model

First off, straight-line code.

New idea: put preconditions on events, e.g.

$$r = 1 \mid W z 2$$

is an execution of ($z := r + 1$).

The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

Model

Attacks

Experiments

Conclusions

Compositional pomset model

The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

First off, straight-line code.

New idea: put preconditions on events, e.g.

$$\boxed{W y 1} \quad \boxed{r = 1 \mid W z 2}$$

is an execution of ($y := 1; z := r + 1$).

Note: no dependency because r does not depend on $y := 1$.

Introduction

Model

Attacks

Experiments

Conclusions

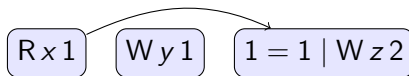
Compositional pomset model

The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

First off, straight-line code.

New idea: put preconditions on events, e.g.



is an execution of $(r := x; y := 1; z := r + 1)$.

Note: dependency because r depends on $r := x$.

Also note: performing a substitution $[1/r]$.

Introduction

Model

Attacks

Experiments

Conclusions


Compositional pomset model

The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

First off, straight-line code.

New idea: put preconditions on events, e.g.



is an execution of $(r := x; y := 1; z := r + 1)$.

Visualize: elide tautologies

Introduction

Model

Attacks

Experiments

Conclusions

Compositional pomset model

Next, conditionals.

The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

Model

Attacks

Experiments

Conclusions

Compositional pomset model

Next, conditionals.

New idea: an execution of $\text{if } M \{ C \} \text{ else } \{ D \}$
comes from an execution of C *and* an execution of D

The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

Model

Attacks

Experiments

Conclusions

Compositional pomset model

Next, conditionals.

New idea: an execution of $\text{if } M \{ C \} \text{ else } \{ D \}$
comes from an execution of C and an execution of D , e.g.

$$r \neq 0 \mid W y 1$$

is an execution of ($y := 1$)
when $r \neq 0$

Compositional pomset model

Next, conditionals.

New idea: an execution of `if $M\{C\}$ else $\{D\}$` comes from an execution of C and an execution of D , e.g.

$$r = 0 \mid Wy2$$

is an execution of ($y := 2$)
when $r = 0$

Compositional pomset model

Next, conditionals.

New idea: an execution of $\text{if } M \{ C \} \text{ else } \{ D \}$ comes from an execution of C and an execution of D , e.g.

$$r \neq 0 \mid W y 1$$

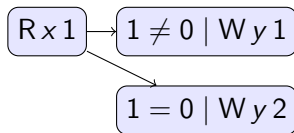
$$r = 0 \mid W y 2$$

is an execution of ($\text{if } (r) \{ y := 1 \} \text{ else } \{ y := 2 \}$)

Compositional pomset model

Next, conditionals.

New idea: an execution of $\text{if } M \{ C \} \text{ else } \{ D \}$ comes from an execution of C and an execution of D , e.g.




is an execution of $(r := x; \text{if } (r) \{ y := 1 \} \text{ else } \{ y := 2 \})$

Compositional pomset model

Next, conditionals.

New idea: an execution of $\text{if } M \{ C \} \text{ else } \{ D \}$ comes from an execution of C and an execution of D , e.g.



is an execution of $(r := x; \text{if } (r) \{ y := 1 \} \text{ else } \{ y := 2 \})$

Visualize: elide tautologies and cross out unsatisfiable

Compositional pomset model

But...

The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

Model

Attacks

Experiments

Conclusions

Compositional pomset model

But... any execution of C should be
an execution of $\text{if } M \{ C \} \text{ else } \{ C \}$

The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

Model

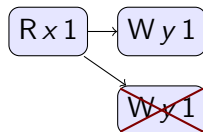
Attacks

Experiments

Conclusions

Compositional pomset model

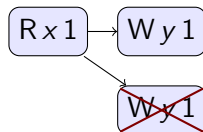
But... any execution of C should be an execution of $\text{if } M \{ C \} \text{ else } \{ C \}$, e.g.



is an execution of $(\text{if } x \{ y := 1 \} \text{ else } \{ y := 1 \})$

Compositional pomset model

But... any execution of C should be an execution of $\text{if } M \{ C \} \text{ else } \{ C \}$, e.g.




is an execution of $(\text{if } x \{ y := 1 \} \text{ else } \{ y := 1 \})$, but so is



Compositional pomset model

But... any execution of C should be an execution of $\text{if } M \{ C \} \text{ else } \{ C \}$, e.g.



is an execution of $(\text{if } x \{ y := 1 \} \text{ else } \{ y := 1 \})$, but so is



New idea: events from different branches can merge.

Compositional pomset model

Lastly, concurrency.

The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

Model

Attacks

Experiments

Conclusions

Compositional pomset model

Lastly, concurrency.

Old idea: match reads with matching writes (à la C11)

The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

Model

Attacks

Experiments

Conclusions

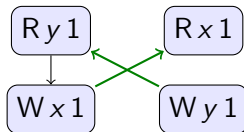
Compositional pomset model

The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Lastly, concurrency.

Old idea: match reads with matching writes (à la C11), e.g.



is an execution of $(x := y \parallel r := x; y := 1)$.

Introduction

Model

Attacks

Experiments

Conclusions

Compositional pomset model

The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Glossed over some details:

- ▶ 3-valued pomsets for negative constraints $d \not\prec e$,
- ▶ sanity conditions on reads-from,
- ▶ precise rules for dependency,
- ▶ variable declaration,
- ▶ ...

All in the paper!

Introduction

Model

Attacks

Experiments

Conclusions

Information flow example

Imagine a SECRET, protected by a run-time security check:

```
if canRead(SECRET) { ... use SECRET ... } else { ... }
```

For attacker code `canRead(SECRET)` is always false

The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

Model

Attacks

Experiments


Conclusions

Information flow example

Imagine a SECRET, protected by a run-time security check:

```
if canRead(SECRET) { ... use SECRET ... } else { ... }
```

For attacker code `canRead(SECRET)` is always false, e.g.



is an execution of


```
if y { if canRead(SECRET) { x := SECRET } else { x := 1 } }.
```

Information flow example

Imagine a SECRET, protected by a run-time security check:

```
if canRead(SECRET) { ... use SECRET ... } else { ... }
```

For attacker code `canRead(SECRET)` is always false, e.g.



is an execution of

```
if y { if canRead(SECRET) { x := SECRET } else { x := 1 } }.
```

Attacker goal: learn if SECRET is 0 or 1.

Modeling Spectre attack

Spectre uses cache timing to discover if a memory location has been touched.

The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

Model

Attacks

Experiments

Conclusions

Modeling Spectre attack

Spectre uses cache timing to discover if a memory location has been touched.

Glossing over a lot of details, this is

```
if touched(x) { ... } else { ... }
```

The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

Model

Attacks

Experiments

Conclusions

Modeling Spectre attack

Spectre uses cache timing to discover if a memory location has been touched.

Glossing over a lot of details, this is

```
if touched(x) { ... } else { ... }
```

Modeled with a new action ($T x$)

The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

Model

Attacks

Experiments

Conclusions


Modeling Spectre attack

Spectre uses cache timing to discover if a memory location has been touched.

Glossing over a lot of details, this is

```
if touched(x) { ... } else { ... }
```

Modeled with a new action (T_x), e.g.



is an execution of $\text{if touched}(x) \{ y := 1 \}$.

The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

Model

Attacks

Experiments

Conclusions


Modeling Spectre attack

Spectre uses cache timing to discover if a memory location has been touched.

Glossing over a lot of details, this is

$$\text{if touched}(x) \{ \dots \} \text{else} \{ \dots \}$$

Modeled with a new action ($T x$), e.g.



is an execution of $\text{if touched}(x) \{ y := 1 \}$.

Require that any event labelled ($T x$) must be preceded by an event labelled ($R x v$) or ($W x v$).

Modeling Spectre attack

The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

A very simplified Spectre attack:

```
if canRead(SECRET) { a[SECRET] := 1 }  
else if touched(a[0]) { x := 0 }  
else if touched(a[1]) { x := 1 }
```

Introduction

Model

Attacks

Experiments

Conclusions

Modeling Spectre attack


The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

A very simplified Spectre attack:

```
if canRead(SECRET) { a[SECRET] := 1 }  
else if touched(a[0]) { x := 0 }  
else if touched(a[1]) { x := 1 }
```

e.g. with execution



Information flow from SECRET to x .

Introduction

Model

Attacks

Experiments

Conclusions

Modeling Prime+Abort attack

Prime+Abort is an information flow attack on Intel's transactional memory. So first model transactions

The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

Model


Attacks

Experiments


Conclusions

Modeling Prime+Abort attack

Prime+Abort is an information flow attack on Intel's transactional memory. So first model transactions, e.g.



and




are executions of `begin; x := x + 1; end`

Modeling Prime+Abort attack


The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely


Prime+Abort is an information flow attack on Intel's transactional memory. So first model transactions, e.g.



and



are executions of `begin; x := x + 1; end`, but *not*



Introduction

Model

Attacks

Experiments

Conclusions


Modeling Prime+Abort attack

Transactions are fine, but not if we add a reason for an abort.

If the attacker knows that every aborted transaction does so because of a read/write or write/write conflict, then in

```
if canRead(SECRET) { a[SECRET] := 1 } ||  
begin; a[1] := 2; loop; end; x := 1
```

the transaction aborts only when SECRET is 1.




Information flow from SECRET to x .

New store reordering attack


An attack on relaxed memory, *discovered from this model.*

```
y := x || if (y == 0) { x := 1 }  
           else if (canRead(SECRET)) { x := SECRET }  
           else { x := 1; z := 1 }
```

If SECRET is 1, there is an execution:



If SECRET is 2, there is no execution:




New dead store elimination attack

Another attack *discovered from this model*.

```
y := x || x := 1;  
  if (canRead(SECRET)) { if (SECRET) { x := 2 } }  
  else { x := 2 }
```

If SECRET is 0, there is an execution:



If SECRET is 1,
there is an execution:




New dead store elimination attack

Another attack *discovered from this model*.

```
y := x || x := 1;  
  if (canRead(SECRET)) { if (SECRET) { x := 2 } }  
  else { x := 2 }
```

If SECRET is 0, there is an execution:



If SECRET is 1, and dead store elimination is performed, there is *no* execution:



Implementing the new attacks

The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Spectre and Prime+Abort are implemented. What about the new attacks?

Introduction

Model

Attacks

Experiments

Conclusions

Implementing the new attacks

The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Spectre and Prime+Abort are implemented. What about the new attacks?

Yes

Introduction

Model

Attacks

Experiments

Conclusions

Implementing the new attacks

The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Spectre and Prime+Abort are implemented. What about the new attacks?

Yes, under unrealistic assumptions:

- ▶ SECRET is a constant known at compile-time,
- ▶ canRead(SECRET) is a run-time check.

Introduction

Model

Attacks

Experiments

Conclusions

Implementing load/store reordering

x86 assembly generated by gcc for the main thread of a variant of the load-store reordering attack:

If SECRET is 0:

```
mov SECRET(%rip), %eax
mov $1, x(%rip)
test %eax, %eax
je label1
mov $0, x(%rip)
```

label1:

```
mov y(%rip), %eax
test %eax, %eax
sete %eax
```

Writes x then reads y,
so can read 1

If SECRET is 1:

```
mov SECRET(%rip), %eax
mov y(%rip), %eax
mov $1, x(%rip)
test %eax, %eax
sete %eax
```

Reads y then writes x,
so cannot read 1

A forwarding thread copies x to y.

The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

Model

Attacks

Experiments

Conclusions

Implementing load/store reordering

The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

To make this attack more likely, introduce a small delay between write of x and read of y , increases probability of round trip.

Experimentally gcc will reorder load/store across 30 straight-line instructions.

Repeat attack to leak multiple bits, and increase probability of success.

Attack is 99.9% accurate at 100Kbps.

Introduction

Model

Attacks

Experiments

Conclusions

Implementing dead store elimination attack

The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

Model

Attacks

Experiments

Conclusions

DSE attack is similar.

Works against clang as well as gcc.

Attack is 99.9% accurate at 400Kbps for clang, and 2Mbps for gcc.

Also in the paper

Details of the model, semantics, etc.

Temporal logic for proving invariants (e.g. no thin-air read).

More examples.

The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

Model

Attacks

Experiments

Conclusions

Contributions

A model of program execution that includes speculation.

Examples including existing information flow attacks on branch prediction and transactional memory, and new attacks on optimizing compilers.

Experimental evidence about how practical it is to mount the new class of attacks

A temporal logic which supports compositional proof.

<https://github.com/chicago-relaxed-memory/spec-eval>

The Code That
Never Ran:
Modeling Attacks
on Speculative
Evaluation

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

Model

Attacks

Experiments

Conclusions