

**PENGUIN
BOOKS**

MYSTERY & CRIME

THE CODE THAT NEVER RAN

MYSTERY & CRIME

CRAIG DISSELKOEN ALAN JEFFREY
RADHA JAGADEESAN JAMES RIELY

COMPLETE



UNABRIDGED

3 January 2018



A day out at the Tate Modern

The Code That Never Ran

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

Spectre

Optimizations

Simplified Spectre

Results

Experiments

Conclusions

3 January 2018



The Code That Never Ran

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

Spectre

Optimizations

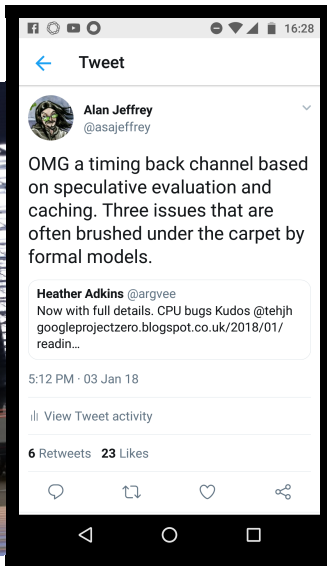
Simplified Spectre

Results

Experiments

Conclusions

3 January 2018



The Code
That Never Ran

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

Spectre

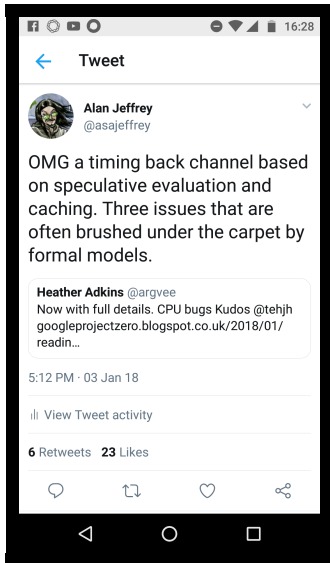
Optimizations

Simplified Spectre

Results

Experiments

Conclusions



Attacks bypass run-time security checks.

Can bypass array bounds checks,
and read whole process memory.

Can be exploited from JS,
so evil.ad.com can read your bank.com data.

Attacks *speculative evaluation*
hardware optimization.

Optimizations in hardware

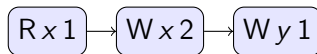
The Code
That Never Ran

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

A lie we tell programmers:
“computers execute instructions one after the other.”

$$x := x + 1; y := 1$$

has execution:



Introduction

Spectre

Optimizations

Simplified Spectre

Results

Experiments

Conclusions

Optimizations in hardware

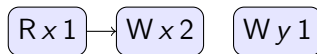
The Code
That Never Ran

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

A lie we tell programmers:
“computers execute instructions one after the other.”

$x := x + 1; y := 1$

has execution where $W y 1$ might happen first:



Introduction

Spectre

Optimizations

Simplified Spectre

Results

Experiments

Conclusions

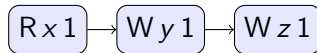
Optimizations in hardware

Another lie we tell programmers:

“only one branch of an if is executed.”

$\text{if}(x) \{ y := 1; z := 1 \} \text{ else } \{ y := 2; z := 1 \}$

has execution:



The Code
That Never Ran

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

Spectre

Optimizations

Simplified Spectre

Results

Experiments

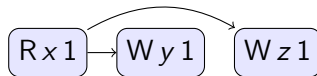
Conclusions

Optimizations in hardware

Another lie we tell programmers:
“only one branch of an if is executed.”

$\text{if}(x) \{ y := 1; z := 1 \} \text{ else } \{ y := 2; z := 1 \}$

has execution where $W z 1$ might happen before $W y 1$:



The Code
That Never Ran

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

Spectre

Optimizations

Simplified Spectre

Results

Experiments

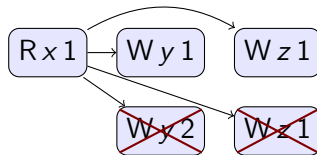
Conclusions

Optimizations in hardware

Another lie we tell programmers:
“only one branch of an if is executed.”

$\text{if}(x) \{ y := 1; z := 1 \} \text{ else } \{ y := 2; z := 1 \}$

has execution where $W y 2$ might happen, then get rolled back:



The Code
That Never Ran

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

Spectre

Optimizations

Simplified Spectre

Results

Experiments

Conclusions

Optimizations in hardware and compilers

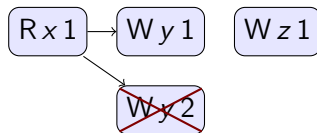
The Code
That Never Ran

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Another lie we tell programmers:
“only one branch of an if is executed.”

$\text{if}(x) \{ y := 1; z := 1 \} \text{ else } \{ y := 2; z := 1 \}$

has execution where $W z 1$ might happen first:



Introduction

Spectre

Optimizations

Simplified Spectre

Results

Experiments

Conclusions

Simplified Spectre

The Code
That Never Ran

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Imagine a SECRET, protected by a run-time security check:

```
if canRead(SECRET) { ... use SECRET ... } else { ... }
```

For attacker code `canRead(SECRET)` is always false

Introduction

Spectre

Optimizations

Simplified Spectre

Results

Experiments

Conclusions

Simplified Spectre

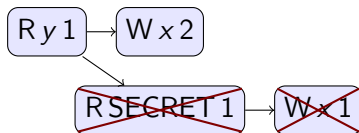
The Code
That Never Ran

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Imagine a SECRET, protected by a run-time security check:

```
if canRead(SECRET) { ... use SECRET ... } else { ... }
```

For attacker code `canRead(SECRET)` is always false, e.g.



is an execution of `if y { if canRead(SECRET) { $x := \text{SECRET}$ } else { $x := 2$ } }`.

Introduction

Spectre

Optimizations

Simplified Spectre

Results

Experiments

Conclusions

Simplified Spectre

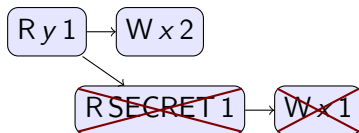
The Code
That Never Ran

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Imagine a SECRET, protected by a run-time security check:

```
if canRead(SECRET) { ... use SECRET ... } else { ... }
```

For attacker code `canRead(SECRET)` is always false, e.g.



is an execution of `if y { if canRead(SECRET) { $x := \text{SECRET}$ } else { $x := 2$ } }`.

Attacker goal: learn if SECRET is 0 or 1.

Introduction

Spectre

Optimizations

Simplified Spectre

Results

Experiments

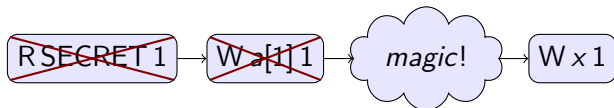
Conclusions

Simplified Spectre

A very simplified Spectre attack:

```
if canRead(SECRET) { a[SECRET]:= 1 }  
else if touched (a[0]) { x:= 0 }  
else if touched (a[1]) { x:= 1 }
```

with execution



Information flow from SECRET to x , *if* there's an implementation of “magic”.

The Code
That Never Ran

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

Spectre

Optimizations

Simplified Spectre

Results

Experiments

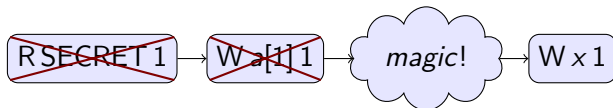
Conclusions

Simplified Spectre

A very simplified Spectre attack:

```
if canRead(SECRET) { a[SECRET]:= 1 }  
else if touched (a[0]) { x:= 0 }  
else if touched (a[1]) { x:= 1 }
```

with execution



Information flow from SECRET to x , *if* there's an implementation of “magic”.

Narrator: there was one.

The Code
That Never Ran

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

Spectre

Optimizations

Simplified Spectre

Results

Experiments

Conclusions

Results

Formalization of pretty pictures as *partially ordered multisets* (Gisher, 1988).

Compositional semantics based on weak memory models (e.g. C11).

Examples modeling Spectre, Spectre mitigations,
PRIME+ABORT attack on transactional memory. . .

The Code
That Never Ran

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

Spectre

Optimizations

Simplified Spectre

Results

Experiments

Conclusions

Results

Formalization of pretty pictures as *partially ordered multisets* (Gisher, 1988).

Compositional semantics based on weak memory models (e.g. C11).

Examples modeling Spectre, Spectre mitigations,
PRIME+ABORT attack on transactional memory...
and a new family of attacks on compiler optimizations.

The Code
That Never Ran

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

Spectre

Optimizations

Simplified Spectre

Results

Experiments

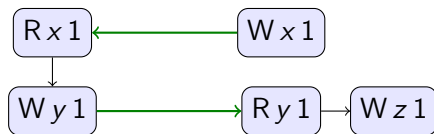
Conclusions

Modeling an attack on compiler optimizations

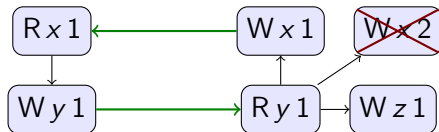
An attacker running two threads (initially $x = y = 0$):

```
y := x || if (y == 0) { x := 1 }  
        else if (canRead(SECRET)) { x := SECRET }  
        else { x := 1; z := 1 }
```

If SECRET is 1, there is an execution:



If SECRET is 2, there is no execution (due to cyclic dependency):



The Code
That Never Ran

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

Spectre

Optimizations

Simplified Spectre

Results

Experiments

Conclusions

Implementing attacks on compiler optimizations

The Code
That Never Ran

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Spectre and Prime+Abort are implemented.

Can we implement the attacks on compiler optimizations?

Introduction

Spectre

Optimizations

Simplified Spectre

Results

Experiments

Conclusions

Implementing attacks on compiler optimizations

The Code
That Never Ran

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Spectre and Prime+Abort are implemented.

Can we implement the attacks on compiler optimizations?

Yes

Introduction

Spectre

Optimizations

Simplified Spectre

Results

Experiments

Conclusions

Implementing attacks on compiler optimizations

The Code
That Never Ran

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

Spectre

Optimizations

Simplified Spectre

Results

Experiments

Conclusions

Spectre and Prime+Abort are implemented.

Can we implement the attacks on compiler optimizations?

Yes, under unrealistic assumptions:

- ▶ SECRET is a constant known at compile-time
- ▶ canRead(SECRET) is a run-time check

Implementing an attack on load/store reordering

Main attacker thread: $x := 1; \text{if}(\text{canRead}(\text{SECRET})) \{ x := \text{SECRET}; \} r := y;$

The Code
That Never Ran

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

Spectre

Optimizations

Simplified Spectre

Results

Experiments

Conclusions

Implementing an attack on load/store reordering

The Code
That Never Ran

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Main attacker thread: $x := 1; \text{if}(\text{canRead}(\text{SECRET})) \{ x := \text{SECRET}; \} r := y;$

When $\text{SECRET} \neq 1$, gcc generates:

```
mov canReadSecret(%rip), %eax
mov $1, x(%rip)
test %eax, %eax
je label1
mov $0, x(%rip)
label1:
mov y(%rip), %eax
```

When $\text{SECRET} = 1$, gcc generates:

```
mov canReadSecret(%rip), %eax
mov y(%rip), %eax
mov $1, x(%rip)
```

Introduction

Spectre

Optimizations

Simplified Spectre

Results

Experiments

Conclusions

Implementing an attack on load/store reordering

The Code
That Never Ran

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Main attacker thread: $x := 1$; if (canRead(SECRET)) { $x := \text{SECRET}$; } $r := y$;

When $\text{SECRET} \neq 1$, gcc generates:

```
mov canReadSecret(%rip), %eax
mov $1, x(%rip)
test %eax, %eax
je label1
mov $0, x(%rip)
label1:
mov y(%rip), %eax
```

Writes x then reads y

When $\text{SECRET} = 1$, gcc generates:

```
mov canReadSecret(%rip), %eax
mov y(%rip), %eax
mov $1, x(%rip)
```

Introduction

Spectre

Optimizations

Simplified Spectre

Results

Experiments

Conclusions

Implementing an attack on load/store reordering

The Code
That Never Ran

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Main attacker thread: $x := 1$; if (canRead(SECRET)) { $x := \text{SECRET}$; } $r := y$;

When $\text{SECRET} \neq 1$, gcc generates:

```
mov canReadSecret(%rip), %eax
mov $1, x(%rip)
test %eax, %eax
je label1
mov $0, x(%rip)
label1:
mov y(%rip), %eax
```

Writes x then reads y

When $\text{SECRET} = 1$, gcc generates:

```
mov canReadSecret(%rip), %eax
mov y(%rip), %eax
mov $1, x(%rip)
```

Conditional has been eliminated!
Reads y then writes x

Introduction
Spectre
Optimizations
Simplified Spectre
Results
Experiments
Conclusions

Implementing an attack on load/store reordering

The Code
That Never Ran

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Main attacker thread: $x := 1$; if (canRead(SECRET)) { $x := \text{SECRET}$; } $r := y$;

When $\text{SECRET} \neq 1$, gcc generates:

```
mov canReadSecret(%rip), %eax
mov $1, x(%rip)
test %eax, %eax
je label1
mov $0, x(%rip)
label1:
mov y(%rip), %eax
```

Writes x then reads y

When $\text{SECRET} = 1$, gcc generates:

```
mov canReadSecret(%rip), %eax
mov y(%rip), %eax
mov $1, x(%rip)
```

Conditional has been eliminated!
Reads y then writes x

Forwarding thread $x := y$ allows attacker to spot the reordering

Introduction
Spectre
Optimizations
Simplified Spectre
Results
Experiments
Conclusions

Implementing an attack on load/store reordering

The Code
That Never Ran

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

Spectre

Optimizations

Simplified Spectre

Results

Experiments

Conclusions

Small delay between write x and read y : increases probability of round trip

gcc will reorder across 30 straight-line instructions

Repeat to leak multiple bits, error correction

Bitwise accuracy 99.99% at 300Kbps

Implementing an attack on dead store elimination

The Code
That Never Ran

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

A similar attack targets dead store elimination

Works on clang + gcc

Bitwise accuracy 99.99% at 1.2Mbps

Introduction

Spectre

Optimizations

Simplified Spectre

Results

Experiments

Conclusions

Contributions

A compositional model of program execution that includes speculation.

Examples including existing information flow attacks on branch prediction and transactional memory, and new attacks on optimizing compilers.

Experimental evidence that the new attacks can be carried out, but only against compile-time secrets.

(Phew, we failed to mount attacks on JIT compilers.)

<https://github.com/chicago-relaxed-memory/spec-eval>

The Code
That Never Ran

Craig Disselkoen,
Radha Jagadeesan,
Alan Jeffrey,
James Riely

Introduction

Spectre

Optimizations

Simplified Spectre

Results

Experiments

Conclusions